

Using Aspera API for Integration and Automation

Unofficial document

Laurent MARTIN

2026/09/29

Contents

1	Introduction	3
1.1	Base rules for a transfer	3
1.2	The basics	4
1.2.1	Client and server roles	4
1.2.2	Server-side storage access (docroot / storage root)	4
1.2.3	Client types and storage capabilities	4
1.2.4	Remote control of clients	5
1.2.5	Transfer authorization	5
1.3	APIs: Components: Server, Client, Application	5
1.4	Addresses and Credentials	6
1.5	Transfer specification, authentication, authorization	6
1.6	Client-side minimum components	6
1.7	First <code>ascp</code> invocation	7
1.8	Sample optional config file for client <code>ascp</code>	7
1.9	CLI	7
1.10	Where to start	8
2	Integration Scenarios	9
2.1	S1- The simplest integration: start a transfer with a script and <code>ascp</code>	9
2.2	S2- Start a transfer with Transfer SDK and listen for events	10
2.3	S3- Start a transfer in a web browser	12
2.4	S4- Start a Server-Server transfer with Node API and SSH credentials	13
2.5	S5- Start a transfer with Token Authorization	15
2.6	S6- Start a Server-Server transfer with Node API and Transfer token	16
2.7	S7- Polling on transfer status on Node API	17
3	Transfer authentication and authorization	19
3.1	Choosing the Right Token Type	19
3.1.1	Architecture comparison	19
3.1.2	ATx Token — Web app as broker	20
3.1.3	Bearer Token — Access control delegated to HSTS	20
3.1.4	Decision guide	21
3.2	A1- SSH Authentication	21
3.3	A2- Aspera Transfer Token Authorization	22
3.3.1	A2: Typical usage	22
3.3.2	A2: Scenario	22
3.3.3	Token generation	22
3.4	A3- Aspera Bearer Token Authorization	23
3.4.1	A3: Typical usage	23
3.4.2	A3: Scenario	24
4	Integration of Aspera Products: Aspera Web Applications	25
4.1	Faspex 5	25
4.1.1	Authentication	25
4.1.2	Transfers (send/receive packages)	25
4.2	IBM Aspera on Cloud	26
4.3	IBM Aspera Shares	26

5	Integration of Aspera Products: Transfer user, Docroot, Storage root, Access keys	27
5.1	Key concepts	27
5.2	Integration of Aspera Products: How to test	27
6	Transfer clients	28
6.1	Native executable: <code>transferd</code>	28
6.1.1	Starting up	28
6.1.2	Examples	28
6.2	Web API	28
6.3	CLI: <code>ascli</code>	28
6.4	Mobile: Android, iOS	28
7	Node API	29
7.1	Overview	29
7.2	What Can You Do With the Node API?	29
7.3	Supported Node Types	29
7.4	Network Interface & Port Configuration	29
7.5	Authentication	30
7.5.1	Basic Authentication	30
7.5.2	Node Users	30
7.5.3	Transfer Node User	30
7.5.4	Master Node User	31
7.5.5	Access Keys	31
7.5.6	Primary Access Keys	31
7.5.7	Sub Access Keys	31
7.5.8	OAuth2 Bearer Tokens	31
7.6	API Endpoint Reference	31
7.6.1	File Management Endpoint Groups	31
7.7	Requests & Responses	32
7.7.1	HTTP Methods	32
7.7.2	Request Format	32
7.7.3	Response Codes	33
7.8	Error Handling	33
7.8.1	Developer Notes	33
7.9	Typical Use Cases	33
7.9.1	Health Check	33
7.9.2	Node Configuration Discovery	33
7.9.3	Access Management	33
7.9.4	File System Operations	34
7.9.5	Transfer Operations	34
7.9.6	Monitoring & Reporting	34

Chapter 1

Introduction

This document examines the various methods for integrating with IBM Aspera APIs to harness their high-performance file transfer capabilities.

IBM Aspera data transfer technology is delivered through a suite of products, each comprising components that expose APIs for integration.

Ultimately, most integrations aim to leverage the IBM Aspera FASP protocol to achieve maximum transfer speeds.

All Aspera transfers involve a client connecting to a server to either push or pull files. The client component is either one of Aspera's client applications, a server, or a custom application built using Aspera's client libraries. The server component is always the IBM Aspera High-Speed Transfer Server (HSTS).




Figure 1.1: Aspera Transfer

1.1 Base rules for a transfer

100% of transfers using Aspera consist of:

- An `ascp` process started in client mode, connecting to a transfer server (SSH or WSS).
- Upon successful connection, an `ascp` process in server mode is started and listens on UDP.
- The connection is direct between client and server using the IP protocol (there can be Aspera proxies in the middle (reverse, forward)), with one TCP connection per transfer session (`ascp` process), and one UDP session.
- Some sort of authorization is required (either just SSH credentials, or a transfer token).
- If the client sends the files, then it is an upload. If the client receives the files, it is a download.

- In all cases, a transfer session is initialized using a "transfer specification", a JSON structure (except direct execution of `ascp` which uses command line options).

1.2 The basics

When using Aspera to transfer data between two storage systems, a few fundamental concepts always apply. These concepts are independent of the specific API, SDK, or product being used (`ascp`, `ascli`, `transferd`, `Node API`, server REST API, etc.).

1.2.1 Client and server roles

An Aspera transfer always involves two distinct roles:

- Client: the side that initiates the transfer.
- Server: the side that accepts the transfer.

The client may either send data to the server (upload) or retrieve data from it (download), but the direction of the data flow does not change the roles: the initiator is always the client.

1.2.2 Server-side storage access (docroot / storage root)

The server side is always configured with access to a storage backend.

This is done through:

- a docroot, or
- an access key + storage root (conceptually equivalent).

The server is responsible for accessing the underlying storage, which can be:

- local filesystem storage, or
- supported object storage (via `PVCL`, such as S3-compatible storage, cloud object stores, etc.).

To do this, the server must be provided with appropriate credentials, which depend on the storage type:

- filesystem permissions for local storage,
- cloud/object-storage credentials for object storage.

Without this configuration, a server cannot read or write data, regardless of the client.

1.2.3 Client types and storage capabilities

The client side can take different forms:

- Simple transfer clients

Examples: `ascp`, `ascli`, `transferd`, `AfD`, custom SDK-based clients.

These clients only support local filesystem storage.

They do not define a docroot or storage root.

Files are always read from or written to the local machine where the client runs.

- Servers acting as clients

An Aspera server can also act as a client.

In that case:

- The server-as-client is configured exactly like a server:
 - ✱ it defines a docroot (or storage root),
 - ✱ it may access object storage,
 - ✱ it requires storage credentials.

- Credentials for object storage are provided either:
 - * directly in the docroot URL, or
 - * via the access key configuration.

This allows transfers between two non-local storage systems, such as object storage to object storage.

1.2.4 Remote control of clients

A client does not need to be manually started by a user. Clients can be remotely controlled:

- A simple client (e.g. `ascli`) can be started remotely via SSH.
- `transferd` exposes a gRPC API that allows remote control of transfers.
- An Aspera server acting as a client can be remotely controlled using its REST API:
 - for example, `POST /ops/transfers` to initiate a transfer.
 - access to this API itself requires proper authentication.

In all cases, the entity triggering the transfer is still the client, even if it is controlled remotely.

1.2.5 Transfer authorization

To initiate a transfer, the client must be authorized to access the server.

This authorization can take different forms, depending on the environment and API:

- SSH-based credentials (user/password or key),
- transfer tokens (Aspera transfer token, JWT, or similar mechanisms).

Without valid authorization, a transfer cannot be started, even if both client and server are correctly configured.

1.3 APIs: Components: Server, Client, Application

These three types of software components provide APIs:

- Transfer Server (HSTS): Node API (REST)
 - Monitor and manage transfers (`GET /ops/transfers`)
 - Start server-to-server transfer (as "transfer client" to remote server) (`POST /ops/transfers`)
 - Get authorization for transfer (typically for web) (`POST /files/*_setup`)
 - Basic file system operations (list files, create folder, etc.) (`GET/POST /files/...`)
 - Supports "watch folder" (includes growing file transfer) (`/watchfolders/...`), `async`, stream, etc.
 - Also provides the server side for FASP transfer (UDP, `ascp`)
- Applications: Faspex, Shares, AoC, Console, Orchestrator API (REST)
 - Provide REST API
 - Do not embed the FASP protocol, but the web app uses the Node API of Transfer Servers
 - Can control and use several Transfer Servers
 - Applications authenticate users, manage users, manage file resources, authorize users for transfers
- Transfer Clients: SDKs
 - Transfer Daemon provides a gRPC interface that can be used from virtually any language.
 - Provides transfer session management (start, monitor)
 - Internally starts one `ascp` process per transfer session.
 - Web SDK, HTTP GW SDK, Desktop SDK: provide the equivalent of Transfer SDK API in the browser (JavaScript)
 - Mobile SDK: the equivalent of Transfer SDK for mobile (Swift for iOS, or Java for Android) (under review)

Note

The Aspera legacy client SDK was called "FaspManager" and is now deprecated (do not use). It provided language-specific implementations: C/C++, Java, Go, Python, .NET, C# (, Ruby). FaspManager2 (based on SWIG) is also deprecated.

1.4 Addresses and Credentials

Examples provided in this document use the following virtual connection information:

- `hsts1.example.com`: Address of HSTS 1.
- `my_hsts1_xfer_user`: A transfer user with SSH credentials on HSTS 1.
- `my_hsts1_xfer_pass`: Password for `my_hsts1_xfer_user`.
- `my_hsts1_node_user`: A node API user (or access key ID).
- `my_hsts1_node_pass`: Password for `my_hsts1_node_user` (or access key secret).
- `hsts2.example.com`: Address of HSTS 2.
- `my_hsts2_node_user`: A node API user (or access key ID).
- `my_hsts2_node_pass`: Password for `my_hsts2_node_user` (or access key secret).

1.5 Transfer specification, authentication, authorization

To start a transfer, the client side must be authorized by the server side. Two types of transfer authorizations are supported (see later):

- SSH authentication/authorization: this is the legacy mode used with Desktop Client or Server-Server transfers.
- Token authorization: this is the mode used, for example, when a web application manages users.

Other parameters are required when starting a transfer, such as transfer direction, while others are optional, such as target speed.

This document will show session parameters using the standardized JSON format "transfer specification" and refer to session parameters as transfer spec. This is the native format used in Web SDKs, `transferd`, Node API and `ascli`. Only `ascp` still does not support transfer spec and still uses regular command line options.

Example of a simple transfer spec for a download using SSH credentials (without token):

```
{
  "remote_host"      : "hsts1.example.com",
  "ssh_port"         : 33001,
  "remote_user"      : "my_hsts1_xfer_user",
  "remote_password"  : "my_hsts1_xfer_pass",
  "direction"        : "receive",
  "destination_root" : ".",
  "paths" : [
    { "source" : "/aspera-test-dir-tiny/200KB.1" }
  ]
}
```

An example with transfer authorization token is provided later.

Use of transfer spec is pervasive in Aspera to describe and start a transfer session (only `ascp` does not use it, yet).

90% of transfer spec parameters are identical between the various APIs (node, `transferd`, Connect, ...).

Parameters can be found in the various API definitions, or with the CLI:

```
ascli conf ascp spec
```

1.6 Client-side minimum components

On the client side, the minimum setup for a FASP transfer consists of two files:

- `ascp` (executable)
- `aspera-license` (free license file, in the same folder as `ascp`, or in `../etc`)

Those files can be found in the [transferd](#) archive.

They can also be found in free clients.

1.7 First ascp invocation

Test `ascp` with a few useful options:

```
ascp -A
```

```
IBM Aspera SDK version 1.1.1.52
ascp version 4.4.2.572 7fd3699
Operating System: MacOSX
FIPS 140-2-validated crypto ready to configure
AES-NI Supported
License max rate=(unlimited), account no.=1, license no.=57999
Enabled settings: desktop_gui
```

```
ascp -h
```

```
ascp -hh
```

```
ascp -DDA
```

```
ascp -DDL-
```

1.8 Sample optional config file for client ascp

On the client side, it is optionally possible to create a configuration file: `aspera.conf` (in the same folder as `ascp`, or in `../etc`).

Example of configuration file: `aspera.conf`, on client side (optional):

```
<?xml version='1.0' encoding='UTF-8'?>
<CONF version="2">
<default>
  <file_system>
    <resume_suffix>.aspera-ckpt</resume_suffix>
    <partial_file_suffix>.partial</partial_file_suffix>
    <replace_illegal_chars>_</replace_illegal_chars>
  </file_system>
</default>
</CONF>
```

Smallest configuration file:

```
<CONF/>
```

Note

This file is mandatory on the server side. On Transfer servers, this file is easily modified with the command `asconfigurator`, which is not available in client applications and SDKs.

1.9 CLI

To test most Aspera APIs and transfer authorization types, one can use the open-source tool `ascli`.

See [manual for installation](#).

`ascli`:

- Connects to all types of Aspera servers
- Provides a logging capability that traces API calls
- Starts transfers using either:
 - Transfer SDK
 - `ascp`
 - Node API
 - Web client (Connect, Aspera for Desktop)
 - HTTP Gateway

Examples in this document are illustrated with `ascli`. The developer can monitor internal API calls, as well as the generated transfer spec, using option `--log-level=trace2`.

1.10 Where to start

There are many scenarios to use Aspera to send files. Get familiar with what an Aspera transfer is. Get the `ascp` binary from transfer SDK or one of the free clients, together with the free license file. See the `transferd` package mentioned earlier, or `ascli`. Execute command line transfers with `ascp` to the demo server using SSH credentials. Do the same with the client method chosen (`transferd` + gRPC). If you send to an Aspera application (Faspex, AoC), use the REST API to create a transfer spec using the application API (and get token authorization).

- Do you need to receive files on your Aspera Server?
- Install an Aspera Transfer Server to test with (evaluation license? dev license?)
- Initiate a transfer to your server.
- Initiate a transfer to your server using the chosen methods (web SDK, node, etc.)

Chapter 2

Integration Scenarios

This section exhibits various file transfer scenarios.

It shows how a FASP transfer is started using a transfer specification.

The transfer is started using various transfer client types/SDKs (CLI, web, app, server).

The transfer specification may contain transfer authorization using either (see later):

- SSH credentials
- Authorization token

Any client type and SDK can be used with any type of transfer authorization: they are independent.

Scenarios contain a typical mix of client SDK and transfer authorization type.

2.1 S1- The simplest integration: start a transfer with a script and ascp




Figure 2.1: script

Scenarios:

- I need to replace `scp` or `sftp` in a script and use my existing SSH credentials (password or key)
- I am provided with bare Aspera Transfer credentials, and I need to transfer to that server using my scripts.

Client	API	Authorization
script	command line	SSH

The simplest and lowest-level integration consists of starting an Aspera transfer as a transfer client with a Transfer Server (upload or download). Ultimately, this consists of starting the `ascp` executable in client mode, locally. Credentials required: SSH credentials provided by the admin of the HSTS server. This can be done easily in a script. For example, using bash:

```
#!/bin/bash
ASPERA_SCP_PASS=my_hsts1_xfer_pass ascp --mode=recv --host=hsts1.example.com -P 33001
↪ --user=my_hsts1_xfer_user aspera-test-dir-tiny/200KB.1 .
```

This starts the `ascp` executable with session parameters on the command line (and an environment variable for the password). `ascp` comes with many session options: a manual can be found in `man ascp` or here: [HSTS Doc](#) → `ascp` command reference. This is suitable for a simple script-based integration, but this solution suffers from the following limitations:

- No easy way to get progress feedback programmatically (progress bar on terminal)
- No automatic retry of failed transfers: the script must implement a loop
- No programmatic API: a process must be spawned with arguments

Equivalent with `ascli`: (includes resume)

```
ascli -N server --url=ssh://hsts1.example.com:33001 --username=my_hsts1_xfer_user
↪ --password=my_hsts1_xfer_pass download aspera-test-dir-tiny/200KB.1
```

Note

Display the generated transfer spec with the additional option `--log-level=debug`.

Equivalent with `asession` (from `aspera-cli` gem), using a transfer spec: (includes progress and resume)

```
asession @json:'{"spec":{"remote_host":"hsts1.example.com","remote_user":"my_hsts1_xfer_user",
↪ "ssh_port":33001,"remote_password":"my_hsts1_xfer_pass","direction":"receive",
↪ "destination_root":".", "paths":[{"source":"/aspera-test-dir-tiny/200KB.1"}]}'
```

2.2 S2- Start a transfer with Transfer SDK and listen for events

Scenarios:

- I need to transfer files at high speed from my home-grown application to a central place using basic OS credentials
- I am provided with bare Aspera Transfer credentials, and I need to transfer to that server using my application written in Java, C++, .NET, Python, etc.

Client	API	Authorization
Custom Client Application	Transfer Daemon gRPC	SSH

If the integration to start a simple transfer job is made in an application written in one of the languages C, C++, C#, Java, Python, Go, Rust, etc., then the best option is to use the newer Transfer Daemon. This component can be freely downloaded here: [Transfer SDK](#). It is basically a daemon, `transferd`, that wraps the `ascp` executable, and it comes with examples.

The basic usage of the API is:

- Initialize the library (or gRPC stub)
- Optionally create an event listener
- Create a transfer spec (with all transfer session parameters)
- Start the transfer job
- Monitor the transfer job (events)

The event listener callback will be called every second with statistics and information about the transfer. The same `ascp` executable and `aspera-license` file from previous step can be used (`transferd` package or free client).




Figure 2.2: transfer daemon

Note that in these two possibilities, the transfer is run locally. Development of a mobile app is similar, with the mobile SDK providing a library to start transfers. The Transfer SDK provides several code examples in several languages.

Other examples: [Laurent's Samples](#)

Equivalent with `ascli`:

```
ascli -N server --url=ssh://hsts1.example.com:33001 --username=my_hsts1_xfer_user
↪ --password=my_hsts1_xfer_pass download aspera-test-dir-tiny/200KB.1 --log-level=debug

...
D, [2019-12-09T13:53:09.541005 #58353] DEBUG -- : mgr is a Asperaalm::Fasp::Local
D, [2019-12-09T13:53:09.541110 #58353] DEBUG -- : ts (json)=
{
  "remote_host": "hsts1.example.com",
  "remote_user": "my_hsts1_xfer_user",
  "ssh_port": 33001,
  "remote_password": "my_hsts1_xfer_pass",
  "direction": "receive",
  "destination_root": ".",
  "paths": [
    {
      "source": "aspera-test-dir-tiny/200KB.1"
    }
  ]
}
D, [2019-12-09T13:53:09.541866 #58353] DEBUG -- : ENV={"ASPERA_SCP_PASS"=>"my_hsts1_xfer_pass"},
↪ ARGS=["-k", "2", "--mode", "recv", "--user", "my_hsts1_xfer_user", "--host", "hsts1.example.com",
↪ "-P", "33001", "--dest64",
↪ "--file-list=/Users/laurent/.aspera/mlia/filelists/c9e8015a-cf45-4deb-9bf4-0e28d96d0f90", "Lg=="]
...
```

Lots of debug information: look for the transfer spec in the logs.

Note

Internally `ascli` uses a "transfer agent" which can be the bare `ascp`, or Transfer Daemon, or other Aspera components.

2.3 S3- Start a transfer in a web browser

Scenarios:

- I need to transfer files at high speed using a web browser, being authorized by some web application (Custom, Faspex, Shares, AoC, ...)

Typically, authentication/authorization is performed in the web app, and a transfer token is used to authorize transfers.

Client	API	Authorization
Web Browser	web SDK	Token

If the transfer must be started by a user in the context of a web browser, then the IBM Aspera JavaScript SDK can be used (for Aspera for Desktop and HTTP GW). It consists of a JavaScript library used similarly to the "Transfer SDK" (build session parameters, start transfer, monitor progress).

Reference (includes sample code): [Aspera API Hub](#) → IBM Aspera JavaScript SDK: <https://developer.ibm.com/apis/catalog/aspera--ibm-aspera-sdk/Introduction>

Example using AfD client with `ascli` (here, not in a browser):

```
ascli -N server --url=ssh://hsts1.example.com:33001 --username=my_hsts1_xfer_user
↪ --password=my_hsts1_xfer_pass download aspera-test-dir-tiny/200KB.1 --transfer=desktop
```




Figure 2.3: web client

The transfer spec remains identical to the previous `ascli` example. Only the transfer agent (`--transfer`) is changed from the local transfer SDK to Aspera for Desktop. Aspera for Desktop is almost exclusively used in a web context: `ascli` is used here as an example only.

2.4 S4- Start a Server-Server transfer with Node API and SSH credentials

Scenarios:

- I need to transfer files between two servers in an automated manner or using REST APIs.
- I own the system used as client side (an HSTS), but not the remote system for which I have only SSH credentials.

Machine-to-machine transfer.

Client	API	Authorization
Aspera Transfer Server	Node API	SSH

This is the method used for automated server-to-server transfers. A transfer can be started remotely (using its REST API) on a server to another remote server. There are two servers here: one starts the transfer as client, and connects to the other one. The managing application uses the "Node API" to control the client side of the transfer. Basically: `POST /ops/transfers` with a JSON payload containing session parameters: transfer spec with SSH credentials of the remote server.

Example of use with `ascli`: (for demonstration of use of API)

```
ascli --url=ssh://hsts1.example.com:33001 --username=my_hsts1_xfer_user --password=my_hsts1_xfer_pass
  --transfer=node --transfer-info=@json: '{"url": "https://hsts2.example.com", "username": "my_hsts2_node_user", "password": "my_hsts2_node_pass"}' --log-level=debug server download
  aspera-test-dir-tiny/200KB.1
```




Figure 2.4: transfer daemon

Note

Option `--transfer=node` tells `ascli` to start the transfer remotely.

In this example, the remote server `hsts2.example.com` is used (as client) to download from server `hsts1.example.com`.

2.5 S5- Start a transfer with Token Authorization




Figure 2.5: transfer daemon

Scenarios:

- I need to transfer files between my custom client app and my custom server app, using my own authentication/authorization.

Client	API	Authorization
Custom Client Application	Transfer SDK(client app) Node API (server app)	Token

In previous examples, transfers were started using SSH credentials and no token, using various client application types: command line, Transfer SDK, browser, server (node) or even mobile. Using SSH credentials, the authentication and transfer authorization are provided by the bare Aspera Server and its host operating system. In fact, in many integration cases, a custom (e.g. web) application server takes care of user authentication and transfer authorization (e.g. RBAC). In this case, once the server application has authenticated and authorized a transfer, it will convey this transfer authorization by generating a "transfer authorization token". Then, the transfer client provides this token to the Aspera server to get authorized to transfer to it. Several types are supported:

- Aspera Auth Token (ATx)

- Basic Authorization
- Bearer Token

Token generation is detailed in the next section. The Aspera token type is a simple choice, as no authorization needs to be set up in advance, and the server application generates it on the fly after it has checked that the application-level user effectively has rights. Aspera tokens are encrypted (using the secret `token_encryption_key`, or a dynamic key) and contain the transfer username, transfer direction, file information, and validity period. This token is generated by calling the node API. The validity period is specified in `aspera.conf` (`token_life_seconds`): 1 day by default. Basic and Bearer tokens are used in specific cases (e.g. AoC uses "access keys", where Bearer tokens can be used and validated by a private/public key pair).

The "AsperaClient" uses one of the SDKs for transfer:

- Transfer (gRPC)
- Web (JavaScript)
- Mobile (Swift, Java)
- Node (REST)

2.6 S6- Start a Server-Server transfer with Node API and Transfer token




Figure 2.6: transfer daemon

Scenarios:

- I own both systems used for the transfer (both are HSTS), and I use the node API to get transfer authorization on the destination, and the node API to initiate the transfer on the source system.

Machine-to-machine transfer.

Client	API	Authorization
Aspera Transfer Server	Node API	Token

This is the method used for automated server-to-server transfers if one owns both servers and token-based transfer is preferred. First the managing application must retrieve a transfer authorization (as well as transfer details, such as server address) using the Node API on the server side of the transfer: `POST /files/upload_setup` with transfer request, which generates a transfer spec. Then, the managing application uses the Node API to control the client side of the transfer. Basically: `POST /ops/transfers` with transfer spec generated previously.

A common mistake is to call both APIs on the same server: client side. Incidentally, this could work if the transfer username is the same on both sides, and if the `token_encryption_key` is the same on both sides for that transfer user. But in general, it will fail, either because the transfer user is different, or the token cannot be decrypted (different encryption key). So, really, authorization shall first be generated on the server side, and the transfer shall be initiated on the client side.

Example of use with `ascli`: (for demonstration of use of API)

```
ascli --url=https://hsts1.example.com:9092 --username=my_hsts1_node_user
↪ --password=my_hsts1_node_pass --transfer=node --transfer-info=@json: '{"url": "https://hsts2.
↪ example.com", "username": "my_hsts2_node_user", "password": "my_hsts2_node_pass"}' --log-level=debug
↪ node upload aspera-test-dir-tiny/200KB.1
```

Note

Option `--transfer=node` tells `ascli` to start the transfer remotely, on the node specified in `--transfer-info`, as client. In this example, the remote server `hsts2.example.com` is used (as client) to upload to server `hsts1.example.com`.

2.7 S7- Polling on transfer status on Node API

The Node API endpoint: `GET /ops/transfers` provides information on past and current transfers. Such information is stored on HSTS in the Redis database.

Activation and retention are controlled by `aspera.conf` parameters:

Parameter	Description
<code>activity_retention</code>	Controls how long data is kept in the database. Default: <code>1d 12h</code>
<code>activity_logging</code>	Enables the logging of transfer activity in the database. Default: <code>false</code>
<code>activity_*</code>	Other parameters related to activity logging.

Note

Typically, `activity_retention` can be given a shorter value to keep the database smaller.
`asconfigurator -x "set_server_data;activity_retention,2h"`

By default (no query parameter), it returns information on a number of transfers. Such call returns a response with the `Link` header set:

```
Link: <https://hsts.example.com:9092/ops/transfers?iteration_token=1440571>; rel="next"
```

Calling the URL specified in `Link` will return subsequent results, and so on. To get all information, repeat with the response's `Link` URL until the response is an empty array and the `iteration_token` does not change. When new information is available, a new call will return data and a new `iteration_token`. A number of query parameters allow filtering only

necessary information, for example: `active_only=true`. Using query `view=id` will return only transfer IDs, and specific information can be retrieved using `GET /ops/transfers/<id>`.

For example, using `ascli`:

```
ascli node transfer list --query.active_only=true
```

Example of an algorithm to get notifications of transfer progress (pseudo-code):

```
iteration_token = 1
transfers = []
loop do
  response = GET /ops/transfers?iteration_token=#{iteration_token}
  if response.iteration_token == iteration_token
    # all current responses received: process, then wait
    call process(transfers)
    transfers = []
    sleep 5
  else
    # additional transfers received: store them and try to get more, before processing all
    transfers.concat(response.transfers)
    iteration_token = response.iteration_token
  end
end
```

Chapter 3

Transfer authentication and authorization

Initiating a transfer session in Aspera requires a valid transfer specification (transfer spec), which encapsulates all necessary parameters for the transfer, including authentication and authorization details. This applies universally, regardless of the method used to start the session.

All transfers are executed on the server under a designated transfer user. The method of authorization is decoupled from the session initiation mechanism, meaning any supported authorization method can be used with any startup interface—whether via script, Transfer SDK, Node API, HTTPGW, mobile, or web applications.

Aspera supports multiple authorization mechanisms:

- OS-Level Authentication (SSH-based)
 - Similar to SFTP, this method uses the operating system's user accounts and SSH authentication.
 - Authorization is managed via `aspera.conf`.
 - Commonly used in legacy environments or high-performance computing (HPC) systems where OS-level user control is required.
- Token-Based Authorization
 - Used when user authentication is handled externally (e.g., by a third-party application).
 - Required for modules that do not use SSH authentication, such as:
 - ★ WebSocket Session
 - ★ Legacy HTTP Fallback
 - ★ HTTP Gateway

There are three types of tokens:

- Aspera Transfer Token (ATx)
 - Generated via the Node API (`/files/*_setup`) using the "gen3" method.
- Basic Token
 - Constructed using access key credentials from the Node API.
- Bearer Token
 - OAuth-style token, typically a JWT signed with a private key.
 - Generated by a custom application and verified by the HSTS using the corresponding public key.
 - Permissions are enforced via the Node API (`/permissions`).

If a token is provided and SSH is used for transport, the transfer user must have Aspera SSH public key authentication enabled. In this case, SSH authentication relies on public keys, while the corresponding private keys are embedded in the SDK or client applications.

3.1 Choosing the Right Token Type

The three token types reflect two fundamentally different architectural patterns for access control. Choosing the wrong one is a common source of integration failures.

3.1.1 Architecture comparison

	ATx (Aspera Transfer Token)	Basic Token	Bearer Token
<u>Node User type</u>	Node API User or Access Key	Access Key	Access Key
<u>File addressing</u>	Path-based (Gen3 <code>/files/*</code>)	ID-based or Path-based	ID-based (Gen4 <code>/files/{id}</code>)
<u>Token generated when</u>	Per transfer request	Permanent (AK credentials)	Once, reused during validity period
<u>Token generated by</u>	Web app via Node API	Client app from AK credentials	Client app using <code>/token</code> endpoint of web app (based on app and user private keys)
<u>Access control enforced by</u>	The web app (before calling Node API)	Access Key scope	HSTS directly (via <code>/permissions</code>)
<u>Typical use case</u>	Custom web app acting as transfer broker	Simple programmatic access	AoC integration, advanced permission model
<u>Complexity</u>	Medium	Low	High

3.1.2 ATx Token – Web app as broker


This is the recommended pattern when a custom web application (or an Aspera application such as Faspex or Shares) sits between the end user and the HSTS.

The client application never calls the Node API directly. Instead:

1. The client authenticates with the web app using whatever mechanism the app supports (SSO, OAuth, cookie, etc.).
2. The client requests a transfer through the web app's own API.
3. The web app enforces its own access control logic, then calls the Node API (`POST /files/upload_setup` or `POST /files/download_setup`) to generate a transfer spec containing a signed ATx token.
4. The web app returns that transfer spec to the client.
5. The client passes the transfer spec directly to the transfer engine (`transferd`, Web SDK, etc.) to execute the transfer.

The HSTS only sees a valid token — it does not know anything about the web app's user model. Access control is entirely the web app's responsibility.

This works with both Node API Users and Access Keys on the HSTS side.




3.1.3 Bearer Token – Access control delegated to HSTS

This is the more advanced pattern, used when access control should be managed directly inside the HSTS rather than in a broker application. It is required for Aspera on Cloud (AoC) integration and any integration using the Gen4 file ID API.

Key characteristics:

- Files are addressed by file ID, not by path. The client must first resolve paths to IDs via `GET /files/{id}/files` before transferring.
- Permissions are set in advance on the HSTS using `POST /permissions`, associating user identifiers (and optional group identifiers) with specific file IDs and access levels.
- The client generates a JWT Bearer token signed with a private key. The HSTS validates it using the corresponding public key configured in advance.
- A single token can be generated once and reused for its entire validity period — the client does not need to call the Node API per transfer.



The added complexity is the trade-off for eliminating the broker intermediary and gaining a richer, ID-stable permission model suitable for multi-tenant or cloud storage environments.

3.1.4 Decision guide

Use ATx when:

- You are building or integrating with a web application that acts as the entry point for users.
- Access control lives in your application layer.
- You want to work with file paths rather than file IDs.
- You use Node API Users or Access Keys.

Use Bearer Token when:

- You are integrating with Aspera on Cloud (AoC).
- You need permissions to be enforced directly by the HSTS, without a broker.
- You are using Access Keys and the Gen4 `/files/{id}` API.
- You need a single token valid across multiple transfer requests.

Use Basic Token when:

- You have programmatic access to an Access Key and need a lightweight token without calling `*_setup` or generating bearer tokens.
- You are in a controlled environment where the Access Key itself provides sufficient access scoping.

3.2 A1- SSH Authentication

For simple transfer initiation in the client-server or server-server case, no prior application-provided authorization is required. The username / pass or key has been pre-arranged between the server provider and client. See the transfer spec example earlier. This is often used in machine-to-machine transfers where authentication is based on a technical operating system user on the server side and SSH. Transfer authorization is based on a top-level `docroot` specified in `aspera.conf` on the server.

Typical end-user usage: simple client-server. This is the legacy mode of authentication / authorization, like `scp` or `sftp`. In this mode, authentication/authorization is done directly during the transfer session: the transfer server provides the authentication (with the local SSH server) and access control (with `aspera.conf`). No external web application controls user authentication or file access rights. The authentication uses the SSH server (with username + password or SSH private key). File transfers are limited to the user's "docroot" (`aspera.conf`) and access rights on storage. Only sessions started with SSH support this mode (not WebSocket Session/HTTPS, or HTTPGW). Once SSH authentication is done, typically the user can transfer any file with possible limits in `aspera.conf` and file system rights.

3.3 A2- Aspera Transfer Token Authorization

3.3.1 A2: Typical usage

Broker application, e.g. Web-based control:

- An external (to transfer server) application (Aspera or custom web app) first authenticates a user in web context (possibly with an external auth system, like SSO) and authorizes a transfer (based on the application access right management).
- The application generates a transfer spec with an authorization token for every transfer requested.
- Transfer authorization tokens are generated before starting the transfer, on the transfer server, using the node API:
`POST /files/<upload|download|sync>_setup`

3.3.2 A2: Scenario




Figure 3.1: transfer daemon

1. The user is authenticated in the application (web, API, etc.)
2. The user calls the app's API with whatever security mechanism (OAuth token, cookie, etc.)
3. The app checks the user's rights based on its own logic, and then calls the node API to generate a transfer spec with token
4. The transfer spec is transmitted to the client side (e.g. web browser) which transmits it to the local Aspera client, which eventually connects to the transfer server where it is authorized for transfer.

3.3.3 Token generation

When user authentication and authorization are managed by a third-party application, then the application needs to provide the client side with a secure "token" that will carry the transfer authorization. The application uses the node API endpoints:

- `/files/upload_setup`
- `/files/download_setup`
- `/files/sync_setup`

This returns a `transfer spec` that can be directly provided to the transfer agent to do the transfer. The authorization is based on an authorization token, generated by the node daemon.

Example:

```
curl -s -u my_hsts2_node_user:my_hsts2_node_pass https://hsts2.example.com/files/download_setup -d
  '{"transfer_requests":[{"transfer_request":{"paths":[{"source":"/aspera/aspera-test-dir-
  ↵ large/1GB"}]}]}'
```

```
{ "transfer_specs": [ { "transfer_spec": {
  "paths": [ { "source": "/aspera/aspera-test-dir-large/1GB" } ],
  "source_root": "",
  "destination_root": "",
  "token": "ATM2_ACSPr3W_d35X0qASXde2bPbvadFuMeG15Hoy9M61_zLbGEAAF465hdKlCgiEOh3gjRfgtf_2MTA",
  "direction": "receive",
  "cipher": "aes-128",
  "tags": null,
  "rate_policy_allowed": "fixed",
  "rate_policy": "fair",
  "target_rate_kbps": 100000,
  "min_rate_kbps": 0,
  "remote_host": "hsts2.example.com",
  "remote_user": "xfer",
  "sshfp": "b42dcb718dd3a04c90e160817d40da9af6e8895a",
  "ssh_port": 33001,
  "fasp_port": 33001,
  "http_fallback": true,
  "http_fallback_port": 8443
} } ] }
```

Aspera Transfer tokens start with `AT` and end with `TA`. They are generated using the `gen3` Node API (`/files/*_setup`). They can also be generated using the utility `astokengen`, but this shall be avoided (legacy).

Generation of Aspera Transfer tokens requires the use of a token encryption key (symmetric) which is either configured statically in `aspera.conf` (`token_encryption_key`) or uses `askms` (dynamic keys). This type of token is simply binary metadata encrypted (and then decrypted) with that key and includes:

- Transfer direction (send/receive, as seen from the server)
- Transfer user
- Date validity
- Hash of authorized file list

When a transfer is initiated with a server and contains such a token, that server checks that each of these elements is valid, otherwise the transfer is refused. This means that in order to be authorized to transfer with a transfer server using such token, the token (and transfer spec) must be generated on that same server (or cluster) using the node API.

3.4 A3- Aspera Bearer Token Authorization

3.4.1 A3: Typical usage

Complex application using OAuth2:

- An external (to transfer server) application (Aspera or custom web app) authenticates users and grants access to its OAuth API in web or API context (possibly with an external auth system, like SSO).
- A private key is used in the application to sign bearer tokens, and the public key is placed on the HSTS in the "access key" to be able to check signatures (and authenticity).
- The Bearer token contains information about the user's identification and group membership.
- The application sets permissions on access to folders on HSTS based on user ID or group ID.
- The user gets a single bearer token from the app which grants access to any transfer.

When a transfer starts on HSTS it will check if the resource (folder/file) is granted access to the user and group IDs found in the provided bearer token.

There is no API to generate a full transfer spec directly, but one can use the `gen3` API to generate a base transfer spec on the HSTS, and then change the token in it with the bearer token.

3.4.2 A3: Scenario




Figure 3.2: transfer daemon

1. The user gets a transfer bearer token for a specific HSTS after authentication to the app.
2. The user uses that bearer token to start transfers. There is no need to re-generate a token for each transfer.
3. The user can also use the bearer token to list files on HSTS directly.
4. Permissions to access files on HSTS are managed out of band, beforehand.

An example of how to configure and generate a bearer token is provided in the `ascli` manual.

The scenario to prepare is:

1. Prepare an Aspera HSTS for use with access keys, and create an access key, for example using the API or `ascli`:
<https://github.com/IBM/aspera-cli/blob/main/docs/README.md#creating-an-access-key>
2. Follow: <https://github.com/IBM/aspera-cli/blob/main/docs/README.md#generating-and-using-a-bearer-token>

Chapter 4

Integration of Aspera Products: Aspera Web Applications

For integration with IBM Aspera applications, use their embedded authentication/authorization and APIs: these, in turn, provide the transfer spec (with token) necessary to initiate a transfer.

For more details, refer to [Aspera API Hub](#).

The general flow is the same as seen earlier with the custom client and server.

Aspera Web Apps provide REST APIs. Some are related to transfer authorization, and others are not directly (e.g. list files, create a user, etc.).

4.1 Faspex 5

4.1.1 Authentication

Faspex 5 uses OAuth 2.0 for authentication and authorization: a Bearer token shall be used for REST API calls. For automated operations, the bearer token shall be obtained using a JWT (JSON Web Token) signed with a private key.

The first step for a client application is to get a client ID and secret from Faspex 5. This is done in the Faspex 5 web interface, in the "API Clients" section.

This client ID and secret, together with the private key, are used to obtain a bearer token by calling the Faspex 5 API endpoint `/token`. Refer to the [Faspex 5 documentation](#) for details, or to examples in GitHub.

A user generates, once, a private/public key pair. The public key is configured in Faspex 5 in the user's profile or in the API client configuration. The private key is used to sign the JWT token, which is then sent to the Faspex 5 API endpoint `/token` to obtain a bearer token. Refer to examples in GitHub.

Once the bearer token is generated, it can be used to call the Faspex 5 REST API endpoints.

4.1.2 Transfers (send/receive packages)

Faspex 5 provides a REST API to manage transfers, including sending and receiving packages.

The first step is to create a transfer spec using the Faspex 5 API. This transfer spec is either for sending a new package, in that case one must first create a package, or for receiving an existing package, identified by a unique ID. A user can list received packages, and then select one to download. The Faspex 5 API provides endpoints to create a transfer spec for sending or receiving packages.

The second step is to use the transfer spec to start the transfer using one of the Aspera transfer agents (Transfer Daemon, Node API, ...).

4.2 IBM Aspera on Cloud

Similar to Faspex 5, AoC uses OAuth 2.0 for authentication and authorization, using API client ID and secret and `/token` with JWT and private key.

There is no endpoint to create a transfer spec directly, but a transfer spec can be built relatively easily: refer to the examples in GitHub.

4.3 IBM Aspera Shares

Same API as Node API (internally, Shares calls the node API for transfer spec generation).

Chapter 5

Integration of Aspera Products: Transfer user, Docroot, Storage root, Access keys

5.1 Key concepts

When a transfer (upload/download) session is received on server side, the transfer process is executed as a given OS user. This user is called the "transfer user". The file system access rights of this user apply (not applicable to object storage). When an application accesses the HSTS Node REST API, it can use one of the following types of authorization:

- Basic: Node user (+ password) [docroot]
- Basic: Access Key (+ secret) [storage root]
- Bearer: OAuth 2.0 Bearer token [storage root]

An initial Node User is necessary on the HSTS to create access keys. Later, sub access keys can be created from access keys. AoC provides an API (ATS) to create access keys on SaaS nodes. A node user is always associated (at creation, see `asnodeadmin`) to a transfer user. A transfer user can be assigned with either (exclusive):

- A single "docroot": the node API can be used "as-is", but access keys cannot be created with it
- Or a list of file restrictions: in this case, the node user can be used to create access keys, whose "storage_root" must match one of the restrictions.

Refer to the HSTS manual section: "Docroot vs. File Restriction". In both cases, node API calls and transfers are "jailed" either in the docroot, or in the storage_root of the access key.

5.2 Integration of Aspera Products: How to test

In order to test, one needs:

- Client-side libraries (Transfer SDK, Connect SDK, Mobile SDK) or API definitions (node API, web app API)
- Server-side testing server

Chapter 6

Transfer clients

6.1 Native executable: `transferd`

The transfer daemon `transferd` provides a gRPC interface for managing transfers. A `proto` definition file is included in the `transferd` package.

6.1.1 Starting up

The daemon must be started in order to manage transfers. It's up to the developer to decide how to start the daemon.

For quick development, the daemon can be started manually in a separate terminal.

For server-level applications, the daemon can be started as a system service (`systemd`).

For end-user applications, the daemon can be started as a background process by the main process.

6.1.2 Examples

Examples in various languages are provided as part of the `transferd` SDK, as well as at:

<https://github.com/laurent-martin/aspera-api-examples/tree/main/app>

6.2 Web API

6.3 CLI: `ascli`

6.4 Mobile: Android, iOS

Chapter 7

Node API

7.1 Overview

The IBM Aspera Node API is a RESTful web service that runs directly on an Aspera server (referred to as a *node*). It provides full programmatic control over the Aspera transfer server environment, enabling developers to manage files, control transfers, configure access, and monitor activity — all over a standard HTTP/HTTPS interface with a JSON-encoded API.

The Node API is part of IBM Aspera High-Speed Transfer Server and Endpoint and acts as a daemon process that exposes endpoints for:

- Remote file operations (browse, create, delete, rename, search)
- Initiating and managing high-speed uploads and downloads (via FASP protocol or HTTP fallback)
- Access and permissions management
- Node monitoring and reporting

The Node API is designed to complement other Aspera APIs, enabling seamless integration with the broader Aspera ecosystem.

7.2 What Can You Do With the Node API?

Capability	Description
<u>Access Management</u>	Manage user access to the node's file system and transfer capabilities
<u>File Transfer</u>	Upload and download files using the FASP protocol or HTTP/HTTPS fallback
<u>Transfer Control</u>	Start, stop, resume, and reconfigure transfers
<u>File Management</u>	Create, delete, rename, search, and browse files and directories
<u>Monitoring</u>	Retrieve node info, available space, transfer events, and bandwidth/volume stats

7.3 Supported Node Types

The Node API is available on several types of Aspera servers:

Node Type	Description
<u>IBM Aspera High-Speed Transfer Server (HSTS)</u>	Self-managed Aspera transfer server
<u>IBM Aspera High-Speed Transfer Endpoint (HSTE)</u>	Self-managed limited version of HSTS
<u>Aspera on Cloud (AoC) Transfer Service (ATS)</u>	Aspera's cloud-based SaaS hosted service (IBM-managed HSTS clusters)

7.4 Network Interface & Port Configuration

The Node API daemon listens on both HTTP and HTTPS interfaces:

Protocol	Default Port	Default Activation
HTTPS	9092	Yes
HTTP	9091	No

Note

Port numbers and activation are configurable through `aspera.conf`. If facing the Internet, it is recommended to place a reverse proxy, such as `nginx`, in front of the node API.

7.5 Authentication

The Node API supports the following authentication mechanisms and types:

- Basic Authentication
 - Node Users
 - * Transfer Node Users
 - * Master Node Users
 - Access Keys
 - * Primary Access Keys
 - * Sub Access Keys
- Bearer token

7.5.1 Basic Authentication

Standard username/password credentials.

There are two main types of node credentials: Node Users and Access Keys.

7.5.2 Node Users

Managed via the `asnodeadmin` utility.

Example: creation of a node user with `asnodeadmin`:

```
asnodeadmin -a -u <API username> -p <API password> -x <OS user>
```

Example: List existing node users:

```
asnodeadmin -l
```

Node users are always associated with a Transfer User, which is an OS-level user. It represents the OS user under which transfers are executed when started or authorized using the node API.

There are two types of Node Users: Transfer Node User and Master Node User.

7.5.3 Transfer Node User

Node Users with a docroot (absolute): they can be used to manage transfers directly. The docroot is associated with the Transfer user, in `aspera.conf`, modified with `asconfigurator`. The docroot is defined as a PVCL URL, which can range from a local file system to object storage (e.g. S3).

! Caution

Cannot be used to integrate with Aspera on Cloud.

7.5.4 Master Node User

Node Users without a docroot, but with file restrictions (which can be open, like `*`). The file restrictions are associated with the Transfer user, in `aspera.conf`, modified with `asconfigurator`. It is used to create Access Keys. It cannot be used for transfers. There is usually a single Master Node User, but there can be more if necessary, if distinct restrictions are needed. The Master Node User is usually associated with the transfer (OS) user: `xfer`, but it is not mandatory.

! Caution

If Access Keys created with this Master Node User are to be used with Aspera on Cloud, or Aspera Enterprise Web Apps, then it is mandatory to name the transfer user as `xfer`.

7.5.5 Access Keys

Created and managed through the Node API. A storage root is associated with it, and must comply with the file restriction associated with the Master Node User used to create it.

7.5.6 Primary Access Keys

Those are created using the node API, using a Master Node User.

The storage root is defined as a JSON in the `storage` key of the access key. The storage `type` can be `local` (local file system), or object storage (e.g. S3).

7.5.7 Sub Access Keys

Those are created using the node API, using a Primary Access Key. The storage `type` is `local`, and the `path` is relative to the storage root of the associated Primary Access Key. The `path` cannot be `/`: it must be a sub-folder within the Primary Access Key's storage root.

7.5.8 OAuth2 Bearer Tokens

Token-based authentication used for integration with cloud services like Aspera on Cloud (AoC). Bearer tokens are always associated with an Access Key. Use of the newer `/files/{id}` endpoint group is required. Bearer tokens contain a user "identifier", managed by the upper level managing application, as well as optional group identifiers.

Access to file resources is granted on the Access Key with the `/permissions` endpoint by providing user and group identifiers and access levels.

7.6 API Endpoint Reference

7.6.1 File Management Endpoint Groups

The Node API provides two distinct groups of file management endpoints. Understanding the difference is important when choosing the right approach for your integration.

7.6.1.1 Group 1 – `/files/{id}`

The next-generation, Aspera-recommended file management endpoints. Manage files by their unique file ID. Sometimes referred to as Gen4.

Endpoint	Purpose
GET <code>/files/{id}</code>	Get metadata for a file or directory
GET <code>/files/{id}/files</code>	List contents of a directory

This requires the use of an Access Key.

File IDs are automatically managed by the node daemon providing the node API, and are stored in a Redis Database together with permissions.

Advantages:

- Manage files by ID (not path), making operations more robust
- Supports bearer token authentication (required for AoC integration)
- Provides more powerful glob filtering for search and browse
- Strongly recommended for cloud-based nodes due to significantly better performance

7.6.1.2 Group 2 – /files/*

An earlier set of file management endpoints (legacy, still supported). Manage files by their file path. Sometimes referred to as Gen3.

Endpoint	Purpose
GET /files/browse	Browse a directory by path
POST /files/create	Create a file or directory
POST /files/delete	Delete a file or directory
POST /files/rename	Rename a file or directory
GET /files/search	Search for files by path

Limitations:

- Supports only Node API credentials or access keys — no bearer token support
- /files/browse and /files/search can be very slow on cloud storage nodes

Important

For cloud-based nodes, always prefer /files/{id}/files over /files/browse.

7.7 Requests & Responses

The Node API follows standard REST conventions.

7.7.1 HTTP Methods

The API uses standard HTTP verbs: GET, POST, PUT, DELETE, and others as appropriate per endpoint.

7.7.2 Request Format

All request bodies must be encoded as JSON:

```
POST /ops/transfers HTTP/1.1
Host: <node-host>:9092
Authorization: Basic <base64-credentials>
Content-Type: application/json

{
  "transfer_requests": [
    {
      "transfer_request": {
        "remote_host": "example.com",
        "direction": "send",
        "paths": [{ "source": "/local/file.txt" }]
      }
    }
  ]
}
```

7.7.3 Response Codes

Status Code	Meaning
200 OK	Request succeeded
201 Created	Resource created successfully
204 No Content	Request succeeded, no response body
400 Bad Request	Malformed request
404 Not Found	Resource not found
409 Conflict	State conflict (e.g., duplicate resource)
500 Internal Server Error	Server-side failure
503 Service Unavailable	Service temporarily unavailable

7.8 Error Handling

When a request fails, the server returns a structured JSON error object:

```
{
  "error": {
    "code": 404,
    "reason": "Not Found",
    "user_message": "Transfer activity logging disabled"
  }
}
```

7.8.1 Developer Notes

- Always check the returned JSON for an `error` field when the HTTP status code is not `2xx`.
- Some known errors may cause the API to return an empty JSON object `{}` instead of an error body — your client code should handle this case explicitly.
- Do not rely solely on the HTTP status code for error detection.

```
// Example: defensive error checking in JavaScript
const data = await response.json();
if (!response.ok || data.error || Object.keys(data).length === 0) {
  throw new Error(data?.error?.user_message ?? "Unknown error from Node API");
}
```

7.9 Typical Use Cases

7.9.1 Health Check

Verify the node is responsive before initiating operations.

```
GET /ping
```

7.9.2 Node Configuration Discovery

Retrieve node configuration — important because some nodes require specific key-value pairs in requests.

```
GET /info
```

7.9.3 Access Management

Create access keys and configure permissions for bearer token users.

```
POST /access_keys
POST /permissions
```

7.9.4 File System Operations

Browse and manage files on the node.

```
GET /files/{id}/files      → list directory contents
POST /files/create         → create a file or directory
POST /files/delete         → delete files
POST /files/rename         → rename a file
```

7.9.5 Transfer Operations

Initiate transfers between nodes or stream data.

```
POST /ops/transfers        → start a transfer
GET /ops/transfers/{id}    → monitor transfer progress
GET /ops/transfers/bandwidth → monitor bandwidth usage
```

7.9.6 Monitoring & Reporting

```
GET /events                → node events (transfer start, file creation, permission changes)
GET /usage                 → transfer volume
GET /space                 → available disk space (HSTS only)
```

End of document